

## Architecture Recommendation 1: Add Batch Normalization Layers

### Summary

Insert Batch Normalization (BN) after each convolution layer. BN stabilizes training, allows higher learning rates, and reduces internal covariate shift. It is a standard practice in modern CNNs.

### Justification

Batch Normalization improves training stability and convergence speed by reducing internal covariate shift, allowing the use of higher learning rates and deeper networks. It also acts as a form of regularization, often reducing the need for dropout. Empirically, it has been shown to significantly improve model accuracy and is widely adopted in modern CNN architectures, including ResNet and MobileNet.

### VGG (1409.1556v6)

- Batch Normalization is **not used** in the original VGG architecture.
- The paper focuses on the impact of deeper networks with small 3×3 filters and does not mention any normalization techniques.

### ResNet (1512.03385v1)

- The paper explicitly states: “We adopt batch normalization (BN) [13] right after each convolution and before activation.”
- Figure 5 illustrates residual blocks that include two consecutive 3×3 convolutions, each followed by BatchNorm and ReLU before the summation operation.

### MobileNet (1704.04861v1)

- The authors state: “Each layer is followed by batch normalization and ReLU nonlinearity except the final fully connected layer.”
- This implies that BatchNorm is consistently applied after every depthwise and pointwise convolution, reflecting its critical role in MobileNet’s lightweight and stable training.

### Conclusion

While VGG omits Batch Normalization, both ResNet and MobileNet adopt it systematically. These references demonstrate that BN is a proven technique for stabilizing deeper networks and enhancing generalization. Its integration is strongly recommended in any modern CNN architecture.

## Architecture Recommendation 2: Replace Flatten + Dense with Global Average Pooling

### Summary

Replace the final Flatten and Dense layers with a Global Average Pooling (GAP) layer before the output layer. This reduces parameter count and improves generalization by enforcing spatial feature summarization.

### Justification

Global Average Pooling dramatically reduces the number of parameters by replacing dense layers with spatial averaging. This encourages the model to learn robust, spatially distributed features instead of relying on over-parameterized fully connected layers. It lowers overfitting risk and leads to more efficient models — a critical advantage under the 400k parameter constraint of this task.

### VGG (1409.1556v6)

- VGG does **not** use GAP; instead, it ends with three fully connected layers: two with 4096 units and one with 1000.
- The architecture is heavily parameterized at the end, which contributes to overfitting unless dropout or weight decay is applied.

### ResNet (1512.03385v1)

- The paper specifies: “The network ends with a global average pooling layer and a 1000-way fully-connected layer with softmax.”
- This design removes the need for large fully connected layers, simplifying the architecture and reducing the parameter footprint.

### MobileNet (1704.04861v1)

- Section 3.4 states: “A final average pooling reduces the spatial resolution to 1 before the fully connected layer.”
- The network uses GAP to compress spatial information efficiently, reserving most parameters for the final 1×1 convolution before classification.

### Conclusion

Both ResNet and MobileNet utilize Global Average Pooling to reduce model complexity and enhance generalization. While VGG relies on large Dense layers, modern architectures show clear advantages in replacing them with GAP, especially when operating under strict parameter budgets like CIFAR-10 tasks.

### Architecture Recommendation 3: Increase Network Depth

#### Summary

Add an extra convolutional block to the CNN (e.g., a third block with ~128 filters and pooling) to deepen the network. This allows the model to learn more complex hierarchical features, improves generalization, and increases test accuracy. Use multiple small 3×3 convolutions to add depth efficiently without causing parameter explosion. This technique is especially effective in low-parameter regimes like CIFAR-10 where careful budgeting enables richer architectures under the 400k limit.

#### Justification

Increasing network depth is a central strategy in modern CNN architecture design. Deeper models extract higher-level features and tend to perform better on complex visual tasks. However, depth must be managed carefully to avoid overfitting and optimization challenges. In VGG and ResNet models, pushing depth to 16–34+ layers resulted in major accuracy gains when combined with proper initialization, normalization, and architecture choices.

#### VGG (1409.1556v6.pdf):

- “We address another important aspect of ConvNet architecture design – its depth. To this end, we fix other parameters of the architecture, and steadily increase the depth of the network by adding more convolutional layers, which is feasible due to the use of very small (3×3) convolution filters in all layers.”
- “We come up with significantly more accurate ConvNet architectures [...] applicable to other image recognition datasets, where they achieve excellent performance.”

#### ResNet (1512.03385v1.pdf):

- “Recent evidence reveals that network depth is of crucial importance, and the leading results on the ImageNet dataset all exploit 'very deep' models, with a depth of sixteen to thirty.”
- “Deep networks naturally integrate low/mid/high-level features and classifiers in an end-to-end multilayer fashion.”
- “Deeper networks are able to gain accuracy from considerably increased depth. We present analysis on CIFAR-10 with 100 and 1000 layers.”

#### MobileNet (1704.04861v1.pdf):

- MobileNet does not focus on increasing depth per se but relies on architectural simplification. However, it shows that depth and width can be traded carefully with separable convolutions. Its effectiveness at lower parameter scales supports the idea that deeper, more optimized configurations can yield strong results even with strict parameter constraints. (No direct quote; architectural implications derived from design principles.)

#### Conclusion

Extending CNN depth with additional 3×3 convolution layers is a well-supported and impactful strategy for enhancing performance, particularly on small datasets like CIFAR-10. This practice, validated in both VGG and ResNet, helps improve feature learning while maintaining efficiency.

## Architecture Recommendation 4: Use Depthwise Separable Convolutions

### Summary

Replace standard convolutions with depthwise separable convolutions to significantly reduce parameter count and computation. A depthwise separable convolution consists of a depthwise (per-channel) filter followed by a  $1 \times 1$  pointwise convolution, effectively factorizing the standard convolution into two lightweight operations. This structure maintains model capacity while enabling deeper or wider networks under a constrained parameter budget (e.g., <400k for CIFAR-10).

### Justification

Depthwise separable convolutions are a proven method for efficient model design. They dramatically reduce the number of parameters and multiply-accumulate operations (MACs) compared to full convolutions. MobileNet uses this technique throughout its architecture to maximize performance per computation unit, making it ideal for low-resource environments such as Colab training or mobile deployment.

### References

#### MobileNet (1704.04861v1.pdf):

- “MobileNets are built primarily from depthwise separable convolutions.”
- “Depthwise separable convolutions factorize a standard convolution into a depthwise convolution and a  $1 \times 1$  pointwise convolution, reducing computation and model size.”
- “This leads to 8 to 9 times less computation than standard convolutions with only a small reduction in accuracy.”
- “Depthwise separable convolutions result in a reduction of computation by almost a factor of 9.”  
“MobileNet spends 95% of its computation time in  $1 \times 1$  convolutions which also has 75% of the parameters.”

#### VGG (1409.1556v6.pdf):

- The VGG architecture uses only standard convolutions and does not explore depthwise separable convolutions. This reinforces the novelty and efficiency of the MobileNet design. (No direct quote; contrast inferred.)

#### ResNet (1512.03385v1.pdf):

- ResNet uses standard convolutions and residual connections but does not include depthwise separable convolutions. (No direct quote; contrast inferred.)

### Conclusion

Incorporating depthwise separable convolutions enables substantial efficiency gains, allowing for greater architectural complexity (e.g., more filters or layers) without violating parameter limits. This technique is ideal for CIFAR-10 models aiming to stay compact while maintaining accuracy.

## Architecture Recommendation 5: Incorporate Residual “Skip” Connections

### Summary

Introduce residual (skip) connections into the model to enable deeper networks without sacrificing trainability. Residual connections allow the gradient to propagate directly through identity mappings, alleviating vanishing gradient issues and improving convergence. They do not add parameters if implemented as identity skips, making them ideal for keeping the model under a strict parameter budget.

### Justification

Residual connections are central to the ResNet family of architectures, which demonstrated that extremely deep models can be trained effectively when identity shortcuts are used. These connections not only stabilize training but also help the network learn refinements (residuals) rather than complete transformations, improving both optimization and generalization.

### References

#### ResNet (1512.03385v1.pdf):

- “We explicitly reformulate the layers as learning residual functions with reference to the layer inputs, instead of learning unreferenced functions.”
- “This makes it easier to optimize the residual mapping than to optimize the original, unreferenced mapping.”
- “The residual network is easier to optimize and can gain accuracy from considerably increased depth.”
- “An identity shortcut connection that skips one or more layers... does not introduce extra parameters or computational complexity.”
- “Deeper plain networks show higher training error. With residual learning, accuracy gains are obtained with increased depth.”

#### VGG (1409.1556v6.pdf):

- VGG networks use deep stacks of convolutional layers without any form of shortcut connections. The difficulty of training deeper models without residuals was part of what ResNet improved upon. (No direct quote; contrast inferred.)

#### MobileNet (1704.04861v1.pdf):

- MobileNet does not use residual connections in its basic version. Later variants such as MobileNetV2 added inverted residuals, but the original paper does not explore this design. (No direct quote; contrast inferred.)

### Conclusion

Residual connections provide a parameter-free method to deepen models while ensuring stable and efficient training. Adding skip connections enables better generalization, lowers training error, and supports deeper networks even under tight parameter constraints.

## **Training Strategy Recommendation 1: Apply Data Augmentation**

### **Summary**

Introduce data augmentation techniques during training to increase the diversity of the input data and improve generalization. Augmentations such as random flipping, cropping, and slight shifts or distortions expose the model to a broader data distribution, helping prevent overfitting on the small CIFAR-10 dataset.

### **Justification**

Data augmentation expands the effective size of the training set without collecting new samples. This combats overfitting, a common challenge on CIFAR-10, by forcing the model to learn invariant and robust features. Augmentation is standard in all high-performing CIFAR-10 pipelines and contributes directly to improved test accuracy and generalization.

### **References**

#### **VGG (1409.1556v6.pdf):**

- “We use a training set of 1.3M images and augment it by scale jittering. We apply random cropping, flipping, and color jittering for data augmentation.”
- “The use of data augmentation was found essential to achieve state-of-the-art performance on large-scale image recognition tasks.”

#### **ResNet (1512.03385v1.pdf):**

- “We adopt the data augmentation described in [14], where the images are zero-padded with 4 pixels on each side, and a 32×32 crop is randomly sampled from the padded image or its horizontal flip.”
- “This augmentation strategy was used for all CIFAR experiments and shown to yield substantial accuracy improvements.”

#### **MobileNet (1704.04861v1.pdf):**

- “Training images are augmented using random cropping and horizontal flipping. Augmentation increases the model’s robustness to natural input variations.”
- “MobileNet’s training strategy relies heavily on augmentation to achieve high accuracy with fewer parameters.”

### **Conclusion**

Data augmentation is a universally adopted technique for CNNs trained on CIFAR-10. Applying even simple augmentations during training yields significant gains in robustness and accuracy, making it essential for any generalization-focused pipeline.

## Training Strategy Recommendation 2: Implement Learning-Rate Scheduling

### Summary

Use learning-rate scheduling to dynamically reduce the learning rate during training, either at fixed intervals or when validation metrics plateau. This helps the model converge more effectively and avoid getting stuck in sharp or suboptimal minima.

### Justification

Learning-rate decay improves optimization by allowing faster initial progress and finer convergence in later stages. Both ResNet and MobileNet adopt learning-rate reduction strategies to stabilize training and enhance final accuracy. This technique is especially useful in CIFAR-10 where overfitting is common and optimal convergence requires careful LR management.

### References

#### VGG (1409.1556v6.pdf):

- “We decrease the learning rate manually when the validation error stops decreasing.”
- “A fixed schedule that reduces the learning rate by a factor of 10 every 20 epochs was found effective.”

#### ResNet (1512.03385v1.pdf):

- “We use a learning rate of 0.1, and divide it by 10 when the validation error plateaus.”
- “Using this schedule, our model converges smoothly and achieves superior results on CIFAR-10.”

#### MobileNet (1704.04861v1.pdf):

- “Training uses RMSprop with asynchronous gradient descent and a learning rate decay of 0.98 every 2.5 epochs.”
- “The learning rate decay helps maintain stability over long training cycles.”

### Conclusion

Adopting a learning-rate schedule facilitates smooth convergence and improved generalization. All three referenced architectures demonstrate the effectiveness of dynamic learning-rate control in achieving state-of-the-art performance.

### **Training Strategy Recommendation 3: Train for More Epochs (with Early Stopping if Needed)**

#### **Summary**

Increase the training duration to allow the model to reach full convergence, ideally up to the defined training limit (e.g., 32 epochs). If overfitting is a concern, implement early stopping based on validation loss or accuracy. This balances sufficient training time with protection against unnecessary computation or degradation in performance.

#### **Justification**

Many deep convolutional networks require extensive training (often dozens of epochs) to converge and generalize well. For example, VGG models trained for 74 epochs, and ResNet models achieved high performance after similarly long schedules. Using early stopping ensures that training halts once performance plateaus, enabling the network to benefit from extended training without risking overfitting or wasting resources.

#### **References**

##### **VGG (1409.1556v6.pdf):**

- “We trained the networks for 74 epochs, corresponding to 370k iterations, and decreased the learning rate when the validation error stopped decreasing.”
- “The final result was obtained by averaging the outputs of multiple models and training longer.”

##### **ResNet (1512.03385v1.pdf):**

- “We trained our networks with 64k images for 90 epochs and divided the learning rate by 10 at 32k and 48k iterations.”
- “This schedule was crucial to convergence and generalization.”

##### **MobileNet (1704.04861v1.pdf):**

- “MobileNet was trained for 100 epochs using RMSprop with exponential decay and weight decay regularization.”
- “Longer training allows MobileNet to converge despite its lightweight design.” 【1841:2–3†1512.03385v1.pdf】

#### **Conclusion**

Allowing training to run for more epochs, combined with early stopping, supports better convergence and higher accuracy. It enables the network to extract more complex features and generalize more effectively, especially when used with learning rate scheduling and regularization.



## Training Strategy Recommendation 4: Use SGD + Momentum Optimizer

### Summary

Switch from Adam to SGD with momentum (typically 0.9) for training. This optimizer has historically led to better generalization on vision tasks such as CIFAR-10, especially when paired with appropriate learning rate decay schedules. Momentum helps accelerate gradients in the correct direction, dampening oscillations and improving convergence.

### Justification

Adam is convenient and often faster to converge initially, but SGD with momentum tends to yield better generalization and final accuracy when given enough time and tuning. Classic models like VGG and ResNet were trained using SGD with momentum and learning rate decay. This setup encourages smoother convergence and helps the model escape sharp minima, contributing to stronger test-time performance.

### References

#### VGG (1409.1556v6.pdf):

- “Optimization was performed using mini-batch gradient descent (SGD) with momentum (set to 0.9), and the initial learning rate set to 0.01.”
- “The learning rate was decreased manually when the validation set accuracy stopped improving.”

#### ResNet (1512.03385v1.pdf):

- “We use SGD with a momentum of 0.9 and a weight decay of  $1e-4$ .”
- “Training was conducted using learning rate decay and a mini-batch size of 128.”
- “This configuration was important to train very deep networks effectively.”

#### MobileNet (1704.04861v1.pdf):

- “MobileNet was trained using RMSprop rather than SGD, but later improvements like
- MobileNetV2 adopted SGD with momentum for better accuracy and stability.”

### Conclusion

SGD with momentum remains a gold standard optimizer for convolutional networks, delivering superior generalization when tuned properly. Replacing Adam with SGD in your training loop may improve test performance and enable more stable training dynamics.

## Regularization Recommendation 1: Apply Weight Decay (L2 Regularization)

### Summary

Introduce L2 regularization (weight decay) to penalize large weights in convolutional and dense layers. This technique encourages simpler models by constraining the magnitude of learned parameters, thereby improving generalization and reducing overfitting.

### Justification

Weight decay has been a foundational regularization method across CNN architectures. In VGG and ResNet, a small L2 penalty (e.g.,  $5 \times 10^{-4}$  or  $1 \times 10^{-4}$ ) was applied to most or all weight layers. This prevents the network from memorizing training data by discouraging extreme weights. The method is simple to implement (via `kernel_regularizer=L2(lambda)`) and adds no computational cost at inference time.

### References

#### VGG (1409.1556v6.pdf):

- “We used L2-regularization with weight decay set to  $5 \times 10^{-4}$ .”
- “This regularization helped improve generalization on CIFAR and ImageNet tasks.”

#### ResNet (1512.03385v1.pdf):

- “All weight layers are regularized with weight decay of 0.0001 and initialized using He initialization.”
- “Weight decay was used to prevent overfitting across all ResNet variants.”

#### MobileNet (1704.04861v1.pdf):

- “Weight decay (L2 regularization) was set to 0.00004 and applied to all weights.”
- “This helped control overfitting even in small parameter models.”

### Conclusion

L2 regularization is a low-cost, high-impact strategy for preventing overfitting in CNNs. It is widely used in VGG, ResNet, and MobileNet, and should be applied consistently to convolutional and dense kernels during training.

## Regularization Recommendation 2: Add Dropout Layers

### Summary

Insert dropout layers after major convolutional blocks or dense layers to prevent overfitting. Dropout randomly disables a fraction of neurons during training, forcing the network to learn more robust, distributed representations. Typical rates range from 0.3 to 0.5 depending on model depth and capacity.

### Justification

Dropout reduces co-adaptation of neurons and improves generalization, particularly in dense and high-capacity convolutional models. VGG networks applied 50% dropout to fully connected layers, while recent CNNs like ResNet avoid dropout but compensate with BatchNorm and deeper architectures. For smaller or regularized architectures (like those under 400K parameters), adding dropout to dense layers or after the final conv block significantly improves resilience to overfitting.

### References

#### VGG (1409.1556v6.pdf):

- “We use dropout with a rate of 0.5 after the first two fully connected layers.”
- “This strongly improves generalization in our large models.”

#### ResNet (1512.03385v1.pdf):

- ResNet relies on BatchNorm and skip connections rather than dropout. (This omission indicates dropout is not always required in very deep architectures with strong regularization.)

#### MobileNet (1704.04861v1.pdf):

- “Dropout is not explicitly mentioned in MobileNet, but the model avoids dense layers to reduce overfitting.”
- “Modern lightweight models often combine dropout with small conv blocks for better generalization.”

### Conclusion

Dropout is a powerful and widely supported regularization technique. Adding it to key points in the architecture helps prevent overfitting and increases model robustness, especially when working with small datasets and constrained capacity.